

Introductory Econometrics with Recitation: TA Session Week 12

Danny Po-Hsien Kang (康柏賢)

May 12, 2026

Today's agenda

- 1 Data Cleaning
 - `pivot_longer()`
 - `pivot_wider()`
- 2 Logarithmic Regression Models
 - Model Setup
 - Coefficient Interpretation
 - Prediction
- 3 Binary Regression Models
 - Conditional Expectation Function
 - LPM, Probit, and Logit
 - Prediction
 - Pseudo R-squared

Today's Dataset

- Today we will use the following CSV file from the OpenIntro Website:
 - `acs12.csv`: Results from the US Census American Community Survey, 2012.
- Please import the dataset into RStudio.

```
acs <- read.csv("data/acs12.csv")

acs <- acs %>%
  na.omit() %>%
  mutate(married = ifelse(married == "yes", 1, 0)) %>%
  filter(income != 0)
```

Data Cleaning: pivot_longer()

- `pivot_longer()` turns several columns into two columns:
 - one column stores the old column names
 - one column stores the old cell values
- It is useful when the same variable is spread across many columns.
- **Syntax:**

```
data_long <- data_wide %>%  
  pivot_longer(  
    cols = c(col_1, col_2, col_3),  
    names_to = "variable_name",  
    values_to = "value"  
  )
```

Data Cleaning: pivot_longer()

- **Example:**

```
score_wide <- data.frame(  
  id = c(1, 2),  
  quiz1 = c(80, 90),  
  quiz2 = c(85, 95)  
)  
  
score_long <- score_wide %>%  
  pivot_longer(  
    cols = c(quiz1, quiz2),  
    names_to = "quiz",  
    values_to = "score"  
  )
```

Data Cleaning: pivot_longer()

- Each student now has one row for each quiz.
- **Output:**

	student	quiz	score
1	Alice	quiz_1	80
2	Alice	quiz_2	85
3	Bob	quiz_1	90
4	Bob	quiz_2	95

Data Cleaning: pivot_wider()

- pivot_wider() does the opposite of pivot_longer().
- It turns values in two columns into several new columns.
- **Syntax:**

```
data_wide <- data_long %>%  
  pivot_wider(  
    names_from = "variable_name",  
    values_from = "value"  
  )
```

- It is useful when each unit should appear in one row.

Data Cleaning: pivot_wider()

- **Example:**

```
score_wide_again <- score_long %>%  
  pivot_wider(  
    names_from = quiz,  
    values_from = score  
  )
```

- **Output:**

	student	quiz_1	quiz_2
1	Alice	80	85
2	Bob	90	95

- In practice, choose the format that makes later analysis easier.

Logarithmic Regression Models

- Consider a basic linear model:

$$\text{income}_i = \beta_0 + \beta_1 \text{hrs_work}_i + \varepsilon_i.$$

- In this model, a one-hour increase in `hrs_work` is associated with a β_1 dollar increase in predicted income.
- To allow nonlinear relationships, we can use logarithms:

$$\text{income}_i = \beta_0 + \beta_1 \log(\text{hrs_work}_i) + \varepsilon_i,$$

$$\log(\text{income}_i) = \beta_0 + \beta_1 \text{hrs_work}_i + \varepsilon_i,$$

$$\log(\text{income}_i) = \beta_0 + \beta_1 \log(\text{hrs_work}_i) + \varepsilon_i.$$

Logarithmic Regression Models

- First, create the log variables.

```
acs <- acs %>%  
  mutate(  
    lnincome = log(income),  
    lnhrs_work = log(hrs_work)  
  )
```

- Then estimate the four models.

```
model1 <- lm_robust(income ~ hrs_work, data = acs)  
model2 <- lm_robust(income ~ lnhrs_work, data = acs)  
model3 <- lm_robust(lnincome ~ hrs_work, data = acs)  
model4 <- lm_robust(lnincome ~ lnhrs_work, data = acs)
```

Logarithmic Regression Models

• Output:

	Linear	Linear-Log	Log-Linear	Log-Log
(Intercept)	-16905.31 *	-101685.91 ***	7.87 ***	3.74 ***
	(7511.36)	(17372.28)	(0.18)	(0.39)
hrs_work	1641.66 ***		0.06 ***	
	(209.69)		(0.00)	
lnhrs_work		41452.42 ***		1.79 ***
		(4993.46)		(0.11)
R ²	0.12	0.09	0.37	0.42
Adj. R ²	0.12	0.09	0.37	0.42
Num. obs.	745	745	745	745
RMSE	55083.74	55838.30	0.96	0.92
*** p < 0.001; ** p < 0.01; * p < 0.05				

Coefficient Interpretation for Different Models

Model	Change in X	Resulting change in $\hat{Y}$
Linear	1 unit	β_1
Linear-Log	1%	$0.01 \times \beta_1$
Log-Linear	1 unit	$(100 \times \beta_1)\%$
Log-Log	1%	$\beta_1\%$

- The interpretation depends on whether Y , X , or both are in logs.
- For log models, the interpretation is often about percentage changes.

Coefficient Interpretation: Examples

- Suppose the estimated coefficient in the Linear-Log model is 41452.42.

$$1\% \text{ increase in hrs_work} \Rightarrow 0.01 \times 41452.42 = 414.52$$

- Income is predicted to increase by about 414.52 dollars.
- Suppose the estimated coefficient in the Log-Log model is 1.79.

$$1\% \text{ increase in hrs_work} \Rightarrow 1.79\% \text{ increase in predicted income.}$$

Prediction with Logarithmic Models

- Suppose a worker increases working hours from 40 to 42.

```
low_hr <- data.frame(  
  hrs_work = 40,  
  lnhrs_work = log(40)  
)  
  
high_hr <- data.frame(  
  hrs_work = 42,  
  lnhrs_work = log(42)  
)
```

- This is also approximately a 5% increase in working hours.

Prediction with Logarithmic Models

- For the Linear model:

```
predict(model1, high_hr) - predict(model1, low_hr) # 3283.326
```

- For the Linear-Log model:

```
predict(model2, high_hr) - predict(model2, low_hr) # 2022.47
```

- For Log-Linear and Log-Log models, predictions are on the log scale.

```
exp(predict(model3, high_hr)) - exp(predict(model3, low_hr))  
# 3569.997  
exp(predict(model4, high_hr)) - exp(predict(model4, low_hr))  
# 2905.476
```

Conditional Expectation Function

- The conditional expectation function is:

$$E[Y | X] = \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k.$$

- If Y is binary, then:

$$E[Y | X] = P(Y = 1 | X).$$

- In this case, OLS gives the linear probability model (LPM):

$$P(Y = 1 | X) = \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k.$$

Why Not Only Use LPM?

- The LPM is easy to estimate and interpret.
- However, it may predict probabilities smaller than 0 or larger than 1.
- To keep predicted probabilities between 0 and 1, we can use nonlinear models.
- Two common models are:
 - Logit model
 - Probit model

Logit and Probit Models

- The Logit model is:

$$P(Y = 1 \mid X) = \frac{\exp(\beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k)}{1 + \exp(\beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k)}.$$

- The Probit model is:

$$P(Y = 1 \mid X) = \Phi(\beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k),$$

where Φ is the standard normal CDF.

Binary Regression: LPM

- We use `lm_robust()` to estimate the linear probability model.
- **Syntax:**

```
lpm <- lm_robust(  
  married ~ age,  
  data = acs  
)
```

Binary Regression: Probit and Logit

- We use `glm()` with `family = binomial()` to estimate logit and probit models.
- **Syntax:**

```
logit <- glm(  
  married ~ age,  
  data = acs,  
  family = binomial(link = "logit")  
)  
  
probit <- glm(  
  married ~ age,  
  data = acs,  
  family = binomial(link = "probit")  
)
```

Binary Regression: Comparison

- **Model comparison:**

```
screenreg(  
  list(lpm, logit, probit),  
  custom.model.names = c("LPM", "Logit", "Probit")  
)
```

- Coefficient interpretation is different across the three models.

Binary Regression: Comparison

• Output:

	LPM	Logit	Probit
(Intercept)	0.08 (0.05)	-1.93 *** (0.26)	-1.16 *** (0.15)
age	0.01 *** (0.00)	0.05 *** (0.01)	0.03 *** (0.00)
R ²	0.12		
Adj. R ²	0.12		
Num. obs.	745	745	745
RMSE	0.46		
AIC		923.09	925.14
BIC		932.32	934.37
Log Likelihood		-459.55	-460.57
Deviance		919.09	921.14
*** p < 0.001; ** p < 0.01; * p < 0.05			

Binary Regression: Prediction

- Create two new observations.

```
young <- data.frame(age = 25)  
old <- data.frame(age = 65)
```

- For LPM, use `predict()` directly.

```
predict(lpm, young)
```

- For logit and probit models, add `type = "response"` to get predicted probabilities.

```
predict(logit, young, type = "response")  
predict(probit, young, type = "response")
```

Binary Regression: Pseudo R-squared

- For logit and probit models, we usually use pseudo R^2 instead of the usual R^2 .
 - Use `pR2()` from the `pscl` package.

```
library(pscl)
pR2(logit)
```

```
fitting null model for pseudo-r2
      11h      11hNull      G2      McFadden      r2ML      r2CU
-459.55   -506.85    94.61         0.09      0.12      0.16
```

- 11h: log-likelihood of the fitted model.
- 11hNull: log-likelihood of the null model.
- McFadden: McFadden's pseudo R^2 .